

# Designing CI-Flows

*CI for thousands of developer*

November 2, 2020

## Disclaimer:

- ▶ Layman's experience from the trenches
- ▶ AKA:  *Davids opinion considered harmful!*

## Who am I

- ▶ Sigma and Ericsson since 2014
  - ▶ Radio Base Station
- ▶ Feature developer → troubleshooter → development environment → Flow Guardian






# Ericsson RBS



- ▶ Thousands of developers
- ▶ All developing for Radio Base Station
  - ▶ Different sub-organizations, different responsibilities
  - ▶ Sigma, one sub-org  $\approx$  700 developers
- ▶ Gerrit / Git / Jenkins / Jira / (Eiffel)
- ▶ + in-house tools

What is CI for you?

# Basic CI



# Big CI Problems

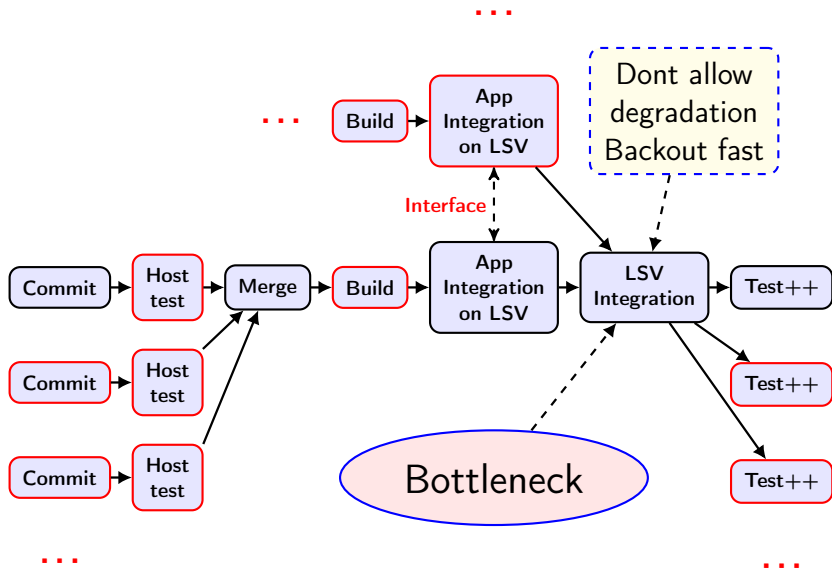
- ▶ Test scope size
  - ▶ Can we run all tests?
  - ▶ Where should tests run?
  - ▶ Are all tests passing?
- ▶ Tracking
  - ▶ Where is my commit?
  - ▶ Is my commit ok?
- ▶ Intermittency
  - ▶ Lots of tests + intermittent tests  $\equiv$  no flow
- ▶ Lead time
  - ▶ Feedback loop
- ▶ Many developers
  - ▶  $\rightarrow$  Many Bottlenecks
  - ▶ Dependencies (expected and unexpected!)



# CI from the CI Flow Plumber's point of view

- ▶ Modularization
- ▶ Logging
- ▶ Non-exhaustive list!
  - ▶ Speed
  - ▶ Stability
  - ▶ Scalability
  - ▶ Reproducibility
  - ▶ ...

# Scalability, a practical example



# Modularization cont

- ▶ One developer/app should not stop flow for all
  - ▶ Bad quality → You don't get to play
  - ▶ Revert/recover first, fix later
- ▶ Needed:
  - ▶ Clean interfaces
  - ▶ Requirements
  - ▶ I.e. good architecture!
- ▶ (Enabler of Agile!)
  - ▶ More defined "sub" responsibilities, better backlogs
  - ▶ Sub-orgs solve similar problems → best solution wins!
- ▶ Shift-left principle
  - ▶ Tests involving only one application tested in app integration
  - ▶ Move tests that often fail earlier in the loops

# Modularization bad things

- ▶ More spread out
  - ▶ Harder to cooperate
  - ▶ Multiple solutions to same problem (alignment)
- ▶ Permissions
  - ▶ "*Why should you have access to my code?*"
- ▶ "Box thinking"
  - ▶ "My box is perfect" → someone elses problem
  - ▶ Remember: All working for same goal

# Big CI Problems

- ▶ Test scope size
  - ▶ Can we run all tests?
  - ▶ Where should tests run?
  - ▶ Are all tests passing?
- ▶ Tracking
  - ▶ Where is my commit?
  - ▶ Is my commit ok?
- ▶ Intermittency
  - ▶ Lots of tests + intermittent tests  $\equiv$  no flow
- ▶ Lead time
  - ▶ Feedback loop
- ▶ Many developers
  - ▶  $\rightarrow$  Many Bottlenecks
  - ▶ Dependencies (expected and unexpected!)

# Way towards Big CI, early actions

1. Parellelism
  - ▶ Enables running many tests
2. Build avoidance / caching
  - ▶ Don't rebuild source/objects that have not changed
  - ▶ Cache objects/build dependencies between consecutive runs
3. Smart testing
  - ▶ Many tests → running all cripples CI
  - ▶ Only run tests that are related to change
4. Invest in Application and CI architecture
  - ▶ Design for testability
    - ▶ Divide application into sub responsibilities (modularization)
    - ▶ Communicate with backwards compatible interfaces
    - ▶ Separation of concerns!
  - ▶ Mocking!

# Running all tests

```
$ cd project-x
$ . ci/setup.sh
$ time apps/app00/test/test.sh
## Running tests for /home/solarus/projects/project-x/apps/app00
# Doing complicated arithmetic (aka sleeping) for 8 seconds ...
# Done!
real 0m8.014s
```

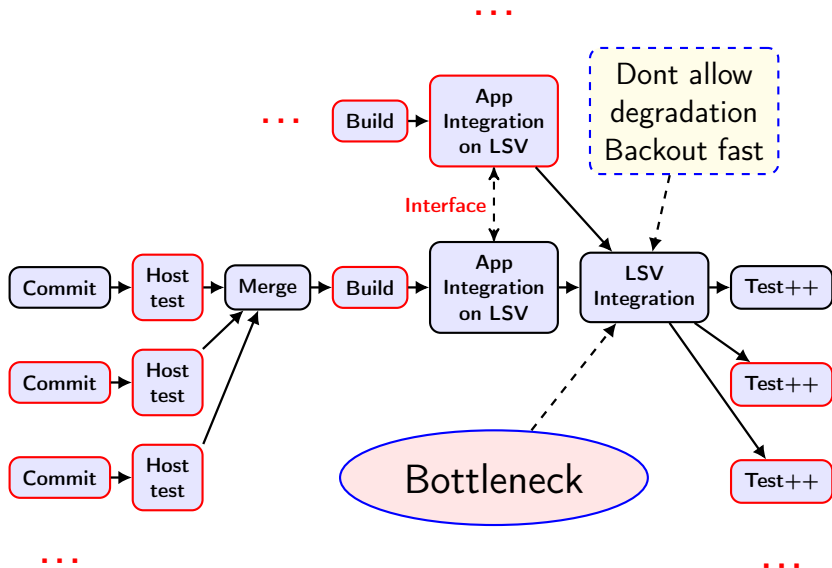
```
$ time find -name test.sh -exec {} \;
## Running tests for /home/solarus/projects/project-x/apps/app04
# Doing complicated arithmetic (aka sleeping) for 0 seconds ...
# Done!
...
## Running tests for /home/solarus/projects/project-x/apps/app03
# Doing complicated arithmetic (aka sleeping) for 28 seconds ...
# Done!
real 11m13.586s
```

# Running all tests

- ▶ In this case 50 suites
  - ▶ Around 15 seconds to finish → on average 12.5 minutes running sequentially
- ▶ Example from one repository:
  - ▶ 1 929 test suites
  - ▶ (1 035 437 lines of test code)
- ▶ Around 15 seconds to finish → about 482 minutes of sequential run time
  - ▶ I.e. a work day...



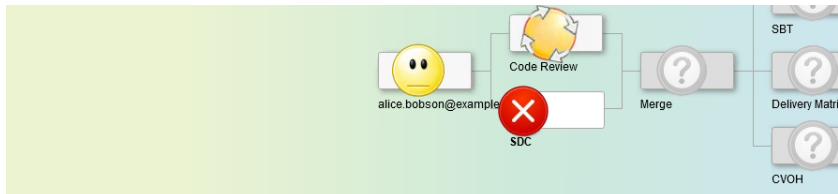
# Scalability, a practical example



# Big CI, a practical example

- ▶ Parallelism
  - ▶ Run as many test as possible at the same time
- ▶ Smart testing
  - ▶ Run only tests related to the application/change
- ▶ Build avoidance
  - ▶ Cache build artifacts that has not changed
  - ▶ ccache
- ▶ Dynamic, developer driven, test scope
  - ▶ Everyone should be able to add/remove test cases from gating guard
  - ▶ Demystifies CI, less scary
  - ▶ Developers are not relying on CI teams for guard update
- ▶ Developer feedback
  - ▶ What has gone wrong
  - ▶ How can I reproduce it?!

# Example of Big CI feedback



## General Information CNA\_SDC (What is this?)

Jenkins job	SDC
Result	Failure
Start Time	Thursday, 29th October 2020, 12:38:14
End Time	Thursday, 29th October 2020, 12:52:04

```
cc.crbs.nc.ucBl coordinatorsSwU.bbIsResHandlerCsSubS
File Edit Format View Help
To reproduce:
git clone .... repo
cd repo
source ci/setup
cd cc/crbs/nc/ucBl/coordinatorsSwU/bbIsR
kp5 -ts coverage
```

## Execute Suites information

[Show all 42 suites](#)

unit	clean	codechecker	coverage	memtest	run	staticCheck
<a href="#">cc/crbs/nc/ucBl/coordinatorsSwU/bbIsResHandlerCsSubSwU/test/bin</a>						
<a href="#">cc/crbs/nc/ucBl/deviceAdaptersSwU/ricmDeviceCtrlCsSubSwU/test/bin</a>						

# Way towards BIG CI, later actions

1. Modularization (architecture)
  - ▶ Integration Testing of subset of complete application
2. Layered testing
  - ▶ E.g. run long running system testing later
3. CI for CI
  - ▶ Run CI changes in same delivery/test flow as applications
4. Logging (big data)
  - ▶ Save test results / data of test case executions!
  - ▶ Impossible for humans to understand complete system
    - ▶ Understanding individual problems is easier
    - ▶ Track where each specific problem occurs in complete CI automatically
5. Stability and Recovery
  - ▶ Not running tests in later loops is expensive!
    - ▶ Long feedback from commit to test execution
    - ▶ If one test fails also following tests → big gap in testing feedback until problem solved
  - ▶ Automatic recovery if application / tests misbehave

# Stability and Recovery

- ▶ Problem: Testing on real hardware → tests or product faults might break test environment
- ▶ Hard to guarantee stand alone testing
- ▶ Solution: Tests try to recover environment to known working configuration before continuing
  - ▶ Never give up!
- ▶ Without recovery, tests after tc failure loose feedback as well
  - ▶ One test failure might lower confidence in big part of system
  - ▶ → Other product/test issues potentially hidden

## Stability: No recovery

```
[=====] Running 5 tests from 1 test suite.  
[-----] Global test environment set-up.  
[-----] 5 tests from ComplicatedApplication  
[ RUN      ] ComplicatedApplication.Test1  
[          OK ] ComplicatedApplication.Test1 (15 min 32s)  
[ RUN      ] ComplicatedApplication.Test2  
[  FAILED    ] ComplicatedApplication.Test2 (10 min 12s)  
[ RUN      ] ComplicatedApplication.Test3  
[  FAILED    ] ComplicatedApplication.Test3 (3 min 5s)  
[ RUN      ] ComplicatedApplication.Test4  
[  FAILED    ] ComplicatedApplication.Test4 (3 min 10s)  
[ RUN      ] ComplicatedApplication.Test5  
[  FAILED    ] ComplicatedApplication.Test5 (3 min 30s)  
[-----] 5 tests from ComplicatedApplication (0 ms total)  
[-----] Global test environment tear-down  
[=====] 5 tests from 1 test suite ran. (0 ms total)
```

## Stability: With recovery

```
[=====] Running 5 tests from 1 test suite.  
[-----] Global test environment set-up.  
[-----] 5 tests from ComplicatedApplication  
[ RUN      ] ComplicatedApplication.Test1  
[          OK ] ComplicatedApplication.Test1 (15 min 12s)  
[ RUN      ] ComplicatedApplication.Test2  
[ FAILED    ] ComplicatedApplication.Test2 (10 min 30s)  
[ RECOVER   ] ComplicatedApplication.Test2 (2 min 3s)  
[ RUN      ] ComplicatedApplication.Test3  
[          OK ] ComplicatedApplication.Test3 (10 min 15s)  
[ RUN      ] ComplicatedApplication.Test4  
[          OK ] ComplicatedApplication.Test4 (5 min 48s)  
[ RUN      ] ComplicatedApplication.Test5  
[          OK ] ComplicatedApplication.Test5 (28 min 13s)  
[-----] 5 tests from ComplicatedApplication (0 ms total)  
[-----] Global test environment tear-down  
[=====] 5 tests from 1 test suite ran. (0 ms total)
```

# Test failure

- ▶ Test case fail → What do you do?



# Logging

- ▶ Remember Big CI Problems:
  - ▶ Many tests+developers+apps/Tracking/Intermittency...
- ▶ Test failed in App Integration →
  - ▶ Test failed before? (same way!)
    - ▶ In same App/other apps?
    - ▶ On certain configurations?
  - ▶ Intermittent?
    - ▶ More intermittent today than last week?

## Logging cont

- ▶ Without data, we are blind to degradations
- ▶ Solution: automatic result tracking!
  - ▶ Test failure messages, configurations, target log analysis
  - ▶ *I.e. store test results and test meta data*
- ▶ Tool to automatically tag known faults in stored results
  - ▶ Tagged faults can be visualized separately (ticket)
  - ▶ → easier to understand
  - ▶ Know if fix helped without reading single test log!

# Summary

- ▶ Parallelism, smart testing, caching crucial
- ▶ Developer driven test scope (dynamic scope), reproducibility
- ▶ Modularization
  - ▶ Test scope per sub-application
  - ▶ Avoids bottlenecks
  - ▶ Good architecture (CI + application)
- ▶ Logging (big data)
  - ▶ Save results for:
    - ▶ Fault tracking
    - ▶ Long term performance tracking
    - ▶ Troubleshooting

# Labs

- ▶ Lab 3: Add smart testing and dynamic test scope for Project-X
- ▶ Lab 4: Set up parallelism for Project-X
- ▶ <https://github.com/dev4242/project-x/tree/main/documents/labs>
  - ▶ Labs, cheat sheet!

Questions?

