

Big CI from scratch

Or how I stopped worrying and started to love the automatic test

March 26, 2021

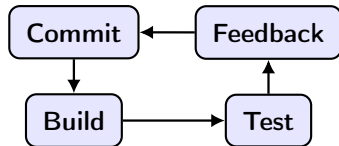
Disclaimer:

- ▶ Layman's experience from the trenches
- ▶ AKA: *Davids opinion considered harmful!*



- ▶ Thousands of developers
- ▶ All developing for Radio Base Station
 - ▶ Different sub-organizations, different responsibilities
 - ▶ Nexer, one sub area \approx couple hundred developers
- ▶ Gerrit / Git / Jenkins / Jira / (Eiffel)
- ▶ + in-house tools

Basic CI



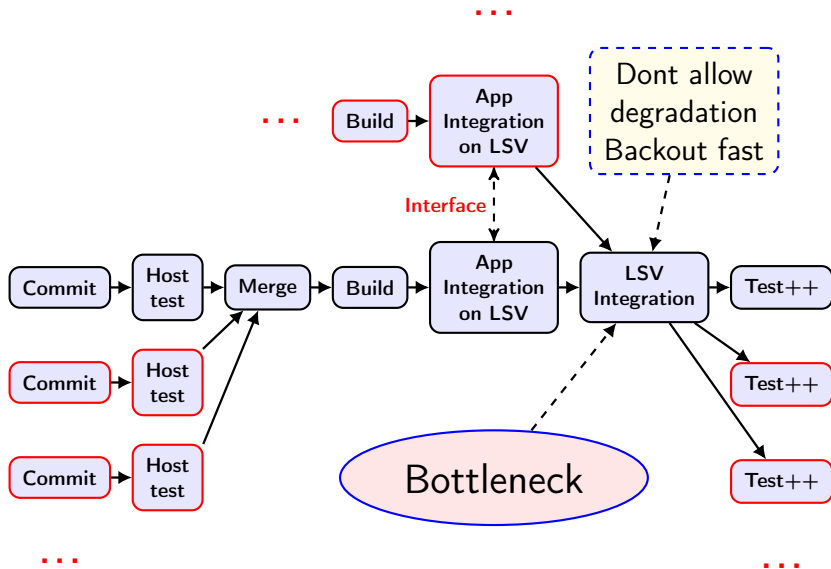
Big CI Problems

- ▶ Test scope
 - ▶ Can we run all tests?
 - ▶ Where should tests run?
 - ▶ Are all tests passing?
- ▶ Tracking
 - ▶ Where is my commit?
 - ▶ Is my commit ok?
- ▶ Intermittency
 - ▶ Lots of tests + intermittent tests \equiv no flow
- ▶ Lead time
 - ▶ Feedback loop
- ▶ Many developers
 - ▶ \rightarrow Many Bottlenecks
 - ▶ Dependencies (expected and unexpected!)

CI from the CI Flow Plumber's point of view

- ▶ Modularization
- ▶ Logging
- ▶ Non-exhaustive list!
 - ▶ Speed
 - ▶ Stability
 - ▶ Reproducibility
 - ▶ Scalability
 - ▶ ...

Modularization



Modularization thinking

- ▶ One developer/app should not stop flow for all
 - ▶ Bad quality → You don't get to play
 - ▶ Revert/recover first, fix later
- ▶ Needed:
 - ▶ Clean interfaces
 - ▶ Requirements
- ▶ (Enabler of Agile!)
 - ▶ More defined "sub" responsibilities, better backlogs
 - ▶ Sub-orgs solve similar problems → best solution wins!

Modularization bad things



Way towards Big CI, early actions

1. Parallelism
 - ▶ Enables running many tests
2. Build avoidance / caching
 - ▶ Don't rebuild source/objects that have not changed
 - ▶ Cache objects/build dependencies between consecutive runs
3. Smart testing
 - ▶ Many tests → running all cripples CI
 - ▶ Only run tests that are related to change
4. Invest in Application and CI architecture
 - ▶ Design for testability
 - ▶ Divide application into sub responsibilities (modularization)
 - ▶ Communicate with backwards compatible interfaces
 - ▶ Separation of concerns!

Running all the tests

```
$ cd project-x
$ . ci/setup.sh
$ time apps/app00/test/test.sh
## Running tests for /home/solarus/projects/project-x/apps/app00
# Doing complicated arithmetic (aka sleeping) for 8 seconds ...
# Done!
real 0m8.014s
```

```
$ time find -name test.sh -exec {} \;
## Running tests for /home/solarus/projects/project-x/apps/app04
# Doing complicated arithmetic (aka sleeping) for 0 seconds ...
# Done!
...
## Running tests for /home/solarus/projects/project-x/apps/app03
# Doing complicated arithmetic (aka sleeping) for 28 seconds ...
# Done!
real 11m13.586s
```

Running all tests

- ▶ In this case 50 suites
 - ▶ Around 15 seconds to finish → on average 12.5 minutes running sequentially
- ▶ Example from one repository:
 - ▶ 1 929 test suites
 - ▶ (1 035 437 lines of test code)
- ▶ Around 15 seconds to finish → about 482 minutes of sequential run time
 - ▶ I.e. a work day...

Questions?



Logging

- ▶ Remember Big CI Problems:
 - ▶ Many tests+developers+apps/Tracking/Intermittency...
- ▶ Test failed in App Integration →
 - ▶ Test failed before? (same way!)
 - ▶ In same App/other apps?
 - ▶ On certain configurations?
 - ▶ Intermittent?
 - ▶ More intermittent today than last week?

Logging cont

- ▶ Without data, we are blind to degradations
- ▶ Solution: automatic result tracking!
 - ▶ Test failure messages, configurations, target log analysis